



MPI Collective Communication

**Irish Centre for High-End Computing
(ICHEC)**
www.ichec.ie



Collective Communication

- Communications involving a group of processes.
- Must be called by all processes in a communicator.
- Examples:
 - Barrier synchronization.
 - Broadcast, scatter, gather.
 - Global sum, global maximum, etc.



Characteristics of Collective Communication

- Optimised Communication routines involving a group of processes
- Collective action over a communicator, i.e. all processes must call the collective routine.
- Synchronization may or may not occur.
- All collective operations are blocking.
- No tags.
- Receive buffers must have exactly the same size as send buffers.




Barrier Synchronization

- C: `int MPI_Barrier(MPI_Comm comm)`
- Fortran: `MPI_BARRIER(COMM, IERROR)`
`INTEGER COMM, IERROR`
- MPI_Barrier is normally never needed:
 - all synchronization is done automatically by the data communication:
 - a process cannot continue before it has the data that it needs.
 - if used for debugging:
 - please guarantee, that it is removed in production.



Broadcast


- C: `int MPI_Bcast(void *buf, int count, MPI_Datatype datatype, int root, MPI_Comm comm)`
- Fortran: `MPI_Bcast(BUF, COUNT, DATATYPE, ROOT, COMM, IERROR)`
`<type> BUF(*)`
`INTEGER COUNT, DATATYPE, ROOT`
`INTEGER COMM, IERROR`



e.g., root=1

- rank of the sending process (i.e., root process)
- must be given identically by all processes


Scatter



- C: int **MPI_Scatter**(void ***sendbuf**, int **sendcount**, MPI_Datatype **sendtype**, void ***recvbuf**, int **recvcount**, MPI_Datatype **recvtype**, int **root**, MPI_Comm **comm**)
- Fortran: **MPI_SCATTER**(**SENDBUF**, **SENDCOUNT**, **SENDTYPE**, **RECVBUF**, **RECVCOUNT**, **RECVTYPE**, **ROOT**, **COMM**, **IERROR**)
 <type> **SENDBUF**(*), **RECVBUF**(*)
 INTEGER **SENDCOUNT**, **SENDTYPE**, **RECVCOUNT**, **RECVTYPE**
 INTEGER **ROOT**, **COMM**, **IERROR**



Gather



- C: `int MPI_Gather(void *sendbuf, int sendcount, MPI_Datatype sendtype, void *recvbuf, int recvcount, MPI_Datatype recvtype, int root, MPI_Comm comm)`
- Fortran: `MPI_GATHER(SENDBUF, SENDCOUNT, SENDTYPE, RECVBUF, RECVCOUNT, RECVTYPE, ROOT, COMM, IERROR)`

<code><type></code>	<code>SENDBUF(*), RECVBUF(*)</code>
<code>INTEGER</code>	<code>SENDCOUNT, SENDTYPE, RECVCOUNT, RECVTYPE</code>
<code>INTEGER</code>	<code>ROOT, COMM, IERROR</code>

```
#include <stdio.h>
#include <mpi.h>
#define SIZE 4

int main(int argc, char **argv){
    int myRank, uniSize, ierror, source;
    float sendbuf[SIZE][SIZE] = {
        {1.0, 2.0, 3.0, 4.0},
        {5.0, 6.0, 7.0, 8.0},
        {9.0, 10.0, 11.0, 12.0},
        {13.0, 14.0, 15.0, 16.0} };
    float recvbuf[SIZE];

    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Comm_rank(MPI_COMM_WORLD,&myRank);
    ierror=MPI_Comm_Size(MPI_COMM_WORLD,&uniSize);
    if (uniSize == SIZE) {
        source = 0;
        ierror=MPI_Scatter(sendbuf,SIZE,MPI_FLOAT,recvbuf,SIZE, MPI_FLOAT,source,MPI_COMM_WORLD);
        printf("rank= %d Results: %f %f %f %f\n", myRank,recvbuf[0],recvbuf[1],recvbuf[2],recvbuf[3]);
    }
    ierror=MPI_Finalize();
    return 0;
}
```

```
rank= 0 Results: 1.0 2.0 3.0 4.0
rank= 1 Results: 5.0 6.0 7.0 8.0
rank= 2 Results: 9.0 10.0 11.0 12.0
rank= 3 Results: 13.0 14.0 15.0 16.0
```



```
program testMPI
use mpi
implicit none

integer :: myRank,uniSize,ierror, source, SIZE
parameter(SIZE=4)
real*4 sendbuf(SIZE,SIZE), recvbuf(SIZE)
data sendbuf /1.0, 2.0, 3.0, 4.0, &
              5.0, 6.0, 7.0, 8.0, &
              9.0, 10.0, 11.0, 12.0, &
              13.0, 14.0, 15.0, 16.0 /

call MPI_Init(ierror)
call MPI_Comm_rank(MPI_COMM_WORLD,myRank,ierror)
call MPI_Comm_Size(MPI_COMM_WORLD,uniSize,ierror)
if (uniSize .eq. SIZE) then
    source = 0
    call MPI_SCATTER(sendbuf,SIZE,MPI_REAL,recvbuf,SIZE,MPI_REAL,source,MPI_COMM_WORLD,ierr)
    print *, 'rank= ',myRank,' Results: ',recvbuf
endif
call MPI_Finalize(ierror)
end program testMPI
```




Global Reduction Operations


- To perform a global reduce operation across all members of a group.
- $d_0 \circ d_1 \circ d_2 \circ d_3 \circ \dots \circ d_{s-2} \circ d_{s-1}$
 - d_i = data in process rank i
 - single variable, or
 - vector
 - \circ = associative operation
 - Example:
 - global sum or product
 - global maximum or minimum
 - global user-defined operation
- floating point rounding may depend on usage of associative law:
 - $[(d_0 \circ d_1) \circ (d_2 \circ d_3)] \circ [\dots \circ (d_{s-2} \circ d_{s-1})]$
 - $(((((d_0 \circ d_1) \circ d_2) \circ d_3) \circ \dots) \circ d_{s-2}) \circ d_{s-1})$

MPI_REDUCE

before MPI_REDUCE




after



- C: `int MPI_Reduce(void *sendbuf, void *recvbuf, int count, MPI_Datatype datatype, MPI_Op op, int root, MPI_Comm comm);`
- Fortran: `MPI_REDUCE(SENDBUF, RECVBUF, COUNT, DATATYPE, OP, ROOT, COMM, IERROR)`

<type> SENDBUF(*), RECVBUF(*)

INTEGER COUNT, DATATYPE, OP, ROOT, COMM, IERROR)



Predefined Reduction Operation Handles

Predefined operation handle	Function
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI LAND	Logical AND
MPI_BAND	Bitwise AND
MPI_LOR	Logical OR
MPI_BOR	Bitwise OR
MPI_LXOR	Logical exclusive OR
MPI_BXOR	Bitwise exclusive OR
MPI_MAXLOC	Maximum and location of the maximum
MPI_MINLOC	Minimum and location of the minimum

User-Defined Reduction Operations

- Operator handles
 - predefined – see table above
 - user-defined
- User-defined function:
 - C: `typedef void MPI_User_function (void *invec, void *inoutvec, int *len, MPI_Datatype *datatype)`
 - Fortran: `FUNCTION USER_FUNCTION (INVEC(*), INOUTVEC(*), LEN, TYPE)`
`<type> INVEC(LEN), INOUTVEC(LEN)`
`INTEGER LEN, TYPE`
- Registering a user-defined reduction function:
 - C: `MPI_Op_create(MPI_User_function *func, int commute, MPI_Op *op)`
 - Fortran: `MPI_OP_CREATE(FUNC, COMMUTE, OP, IERROR)`
- COMMUTE tells the MPI library whether FUNC is commutative.



Variants of Reduction Operations

- **MPI_ALLREDUCE**
 - no root,
 - returns the result in all processes
- **MPI_REDUCE_SCATTER**
 - result vector of the reduction operation is scattered to the processes into the real result buffers
- **MPI_SCAN**
 - prefix reduction
 - result at process with rank i :=
reduction of inbuf-values from rank 0 to rank i