

Introduction to MPI



ICHEC
Irish Centre for High-End Computing



An Roinn Post, Fiontar agus Nuálaíochta
Department of Jobs, Enterprise and Innovation



science foundation ireland
fionduisceacht eolaíochta éireann



AN ROINN
OIDEACHAIS AGUS SCILEANNA
DEPARTMENT OF
EDUCATION AND SKILLS

HEA

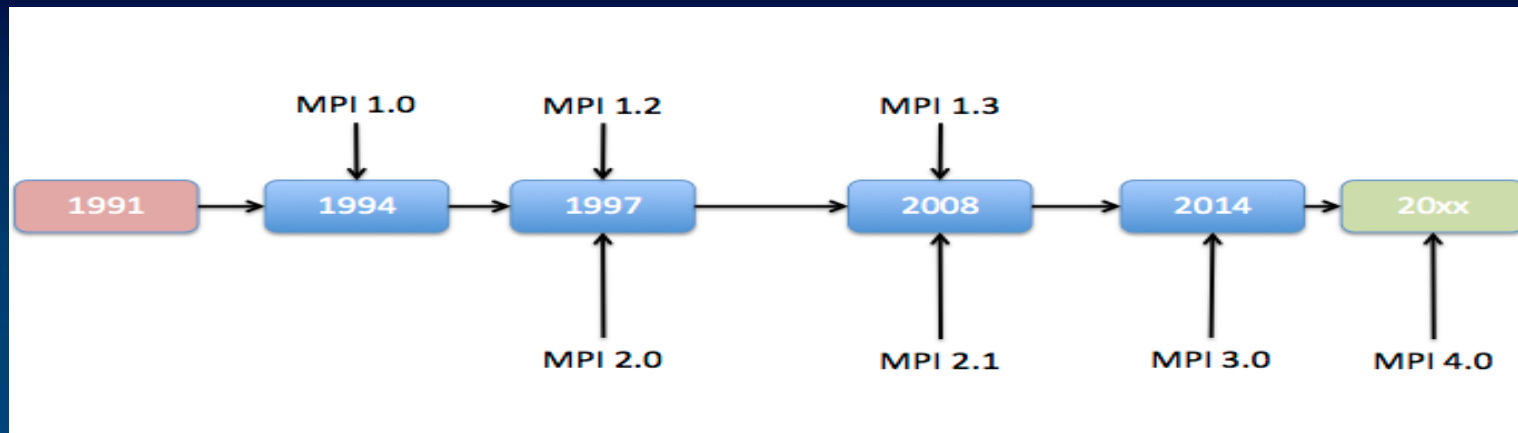
Higher Education Authority
an tUdarás um Ard-Oideachas

www.ichec.ie

MPI (Message Passing Interface)?

- Standardized message passing library specification (IEEE)
 - for parallel computers, clusters and heterogeneous networks
 - not a specific product, compiler specification etc.
 - many implementations, MPICH, LAM, OpenMPI ...
- Portable, with Fortran and C/C++ interfaces.
- Many functions
- Real parallel programming
- Notoriously difficult to debug


Information about MPI



- <http://www.mpi-forum.org/docs/>
- **MPI: The Complete Reference**, Marc Snir and William Gropp et al, The MIT Press, 1998 (2-volume set)
- **Using MPI: Portable Parallel Programming With the Message-Passing Interface and Using MPI-2: Advanced Features of the Message-Passing Interface**. William Gropp, Ewing Lusk and Rajeev Thakur, MIT Press, 1999 – also available in a single volume *ISBN 026257134X*.
- **Parallel Programming with MPI**, Peter S. Pacheco, Morgan Kaufmann Publishers, 1997 – *very good introduction*.
- <https://computing.llnl.gov/tutorials/mpi/>


The Message-Passing Programming Paradigm

- Sequential Programming Paradigm




A processor may run many processes

- Message Passing Programming Paradigm

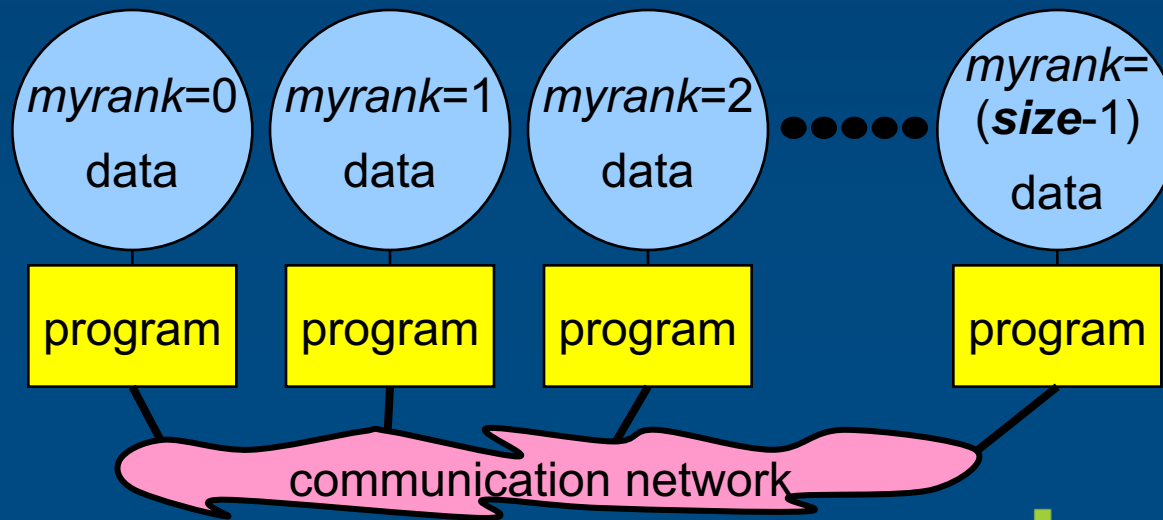


- A **process** is a program performing a task on a **processor/core**
- Each processor/process in a message passing program runs a instance/copy of a **program**:
 - written in a conventional sequential language: C/C++, Fortran, python
 - typically a single program operating on multiple dataset
 - the variables of each sub-program have
 - the same name
 - but different locations (distributed memory) and different data!
 - i.e., all variables are local to a process
 - communicate via special send & receive routines (**message passing**)



Data and Work Distribution

- To communicate together mpi-processes need identifiers:
rank = identifying number
- all distribution decisions are based on the *rank*
 - i.e., which process works on which data



Example : sum of elements of a vector

- **Sequential code**

```
sum = 0
for (int i = 0; i < 1000 ;++i)
    sum = sum + array[i] ;
```

- **parallel code**

```
sum = 0
chunkSize = 1000 / numProc
for (int i= rank*chunkSize; i< (rank + 1)*chunkSize; ++i)
    sum = sum + array[i] ;

if( rank != root)
    send (partialsum) to root
else
    receive(partialsum) from workers
```

On each processor

```
sum = 0
```

```
for (int i = 0; i < 500; ++i)  
    sum = sum + array[i] ;
```

P_0

```
sum = 0
```

```
for (int i = 500; i < 1000; ++i)  
    sum = sum + array[i] ;
```

P_1

- The same program
- The same variables, but different values

What is SPMD

- **Single Program, Multiple Data**
- Same (sub-)program runs on each processor
- MPI allows also MPMD, i.e., **Multiple Program, ...**
 - but some vendors may be restricted to SPMD
 - MPMD can be emulated with SPMD

Emulation of MPMD

- C/C++:

```
main(int argc, char **argv) {
  if (myrank < .... /* process should run the ocean model */) {
    ocean( /* arguments */ );
  } else {
    weather( /* arguments */ );
  }
}
```
- Fortran


```
program forecast
  if (myrank < ... ) then  !! process should run the ocean model
    call ocean ( some arguments )
  ELSE
    call weather ( some arguments )
  endif
end program forecast
```

Message Passing System

- A sub-program needs to be connected to a message passing system
- A message passing system is similar to:
 - phone line
 - mail box
 - fax machine
 - etc.
- MPI:
 - program must be linked with an MPI library
 - program must be started with the MPI startup tool

Message passing

- Messages are packets of data moving between sub-programs
- Necessary information for the message passing system:
 - sending process
 - source location
 - source data type
 - source data size
 - receiving process
 - destination location
 - destination data type
 - destination buffer size




Point-to-Point Communication

- Simplest form of message passing.
- One process sends a message to another.
- Different types of point-to-point communication:
 - synchronous send
 - buffered = asynchronous send


Synchronous Sends

- The sender sends data and waits until it gets an information that the message is received.



Buffered = Asynchronous Sends

- Only know when the message has left.




Blocking Operations

- Some sends/receives may **block** until another process acts:
 - synchronous send operation **blocks until** receive is issued;
 - receive operation **blocks until** message is sent.
- Blocking subroutine returns only when the operation has completed.

Non-Blocking Operations

- Non-blocking operations return immediately and allow the sub-program to perform other work.




Collective Communications

- Collective communication routines are higher level routines.
- Several processes are involved at a time.
- May allow **optimized internal** implementations, e.g., tree based algorithms


Broadcast

- A one-to-many communication.




Reduction Operations

- Combine data from several processes to produce a single result.



Barriers

- Synchronize processes.



Goals and Scope of MPI

- MPI's prime goals
 - To provide a message-passing interface.
 - To provide source-code portability.
 - To allow efficient implementations.
- It also offers:
 - A great deal of functionality.
 - Support for heterogeneous parallel architectures.
- With MPI-2/MPI-3:
 - Important additional functionality.
 - Backward compatibility with MPI-1.