

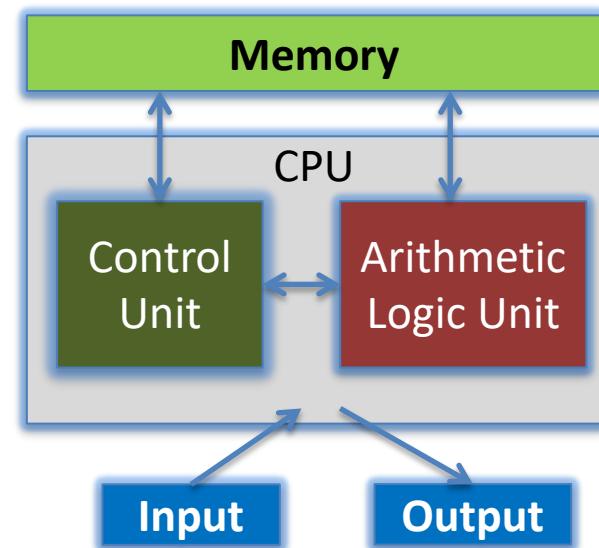


The Fundamentals of Computer & HPC Architectures



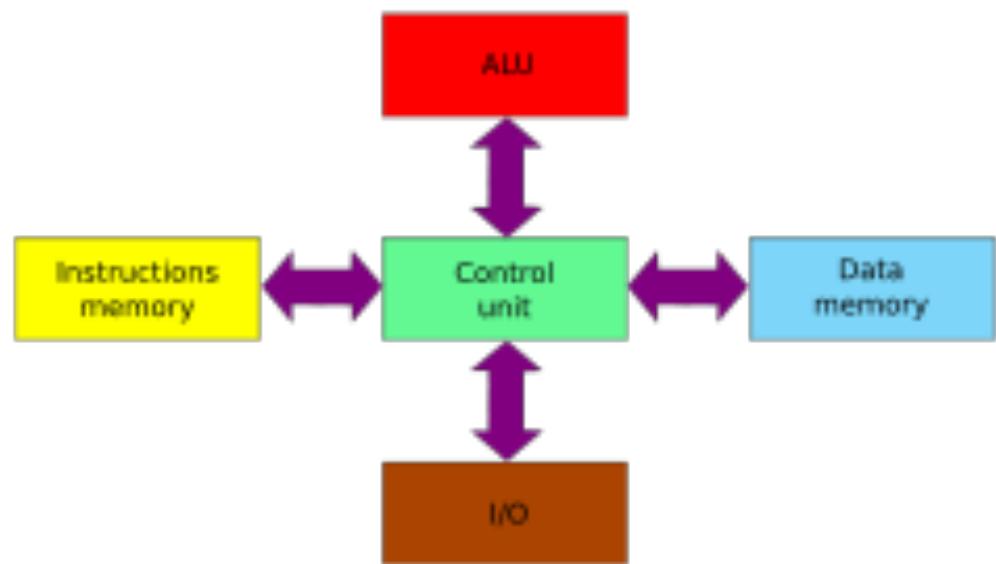
The von Neumann Architecture

- Hungarian-American mathematician John von Neumann proposed the design in 1945
 - Part of the proposal of computer EDVAC, which followed ENIAC
- Four main components:
 - Memory
 - Control Unit
 - Arithmetic Logic Unit
 - Input/Output



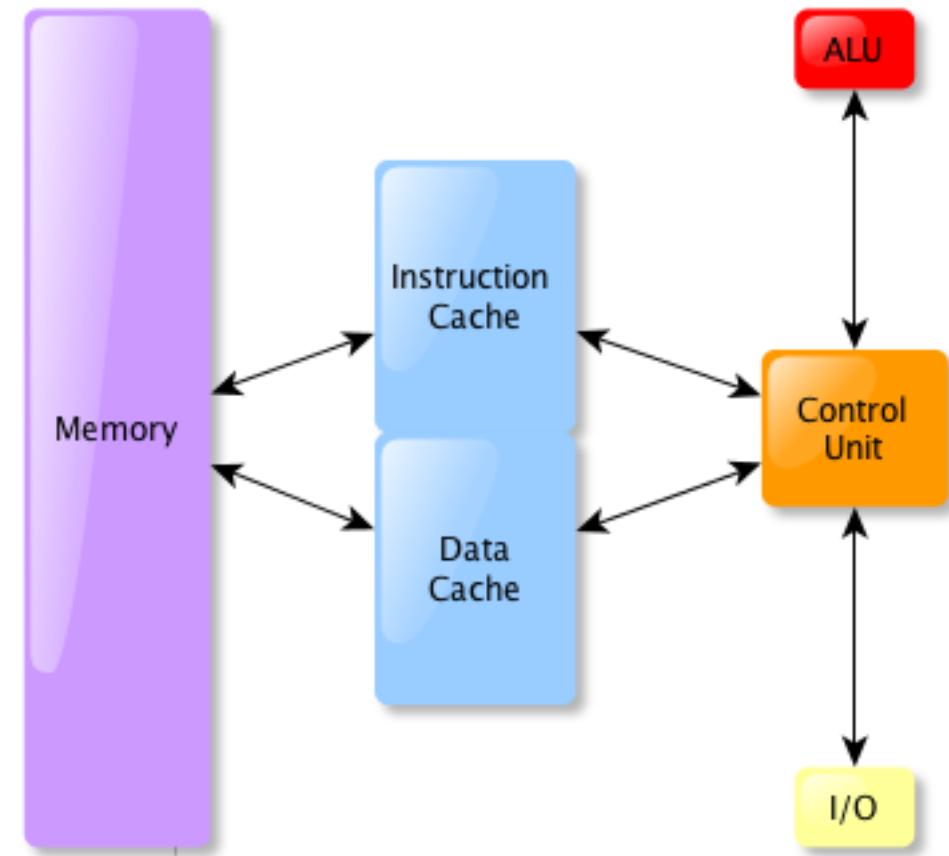
Harvard Architecture

- Instruction read and data memory access can be performed simultaneously
- Different address spaces for program and data memories

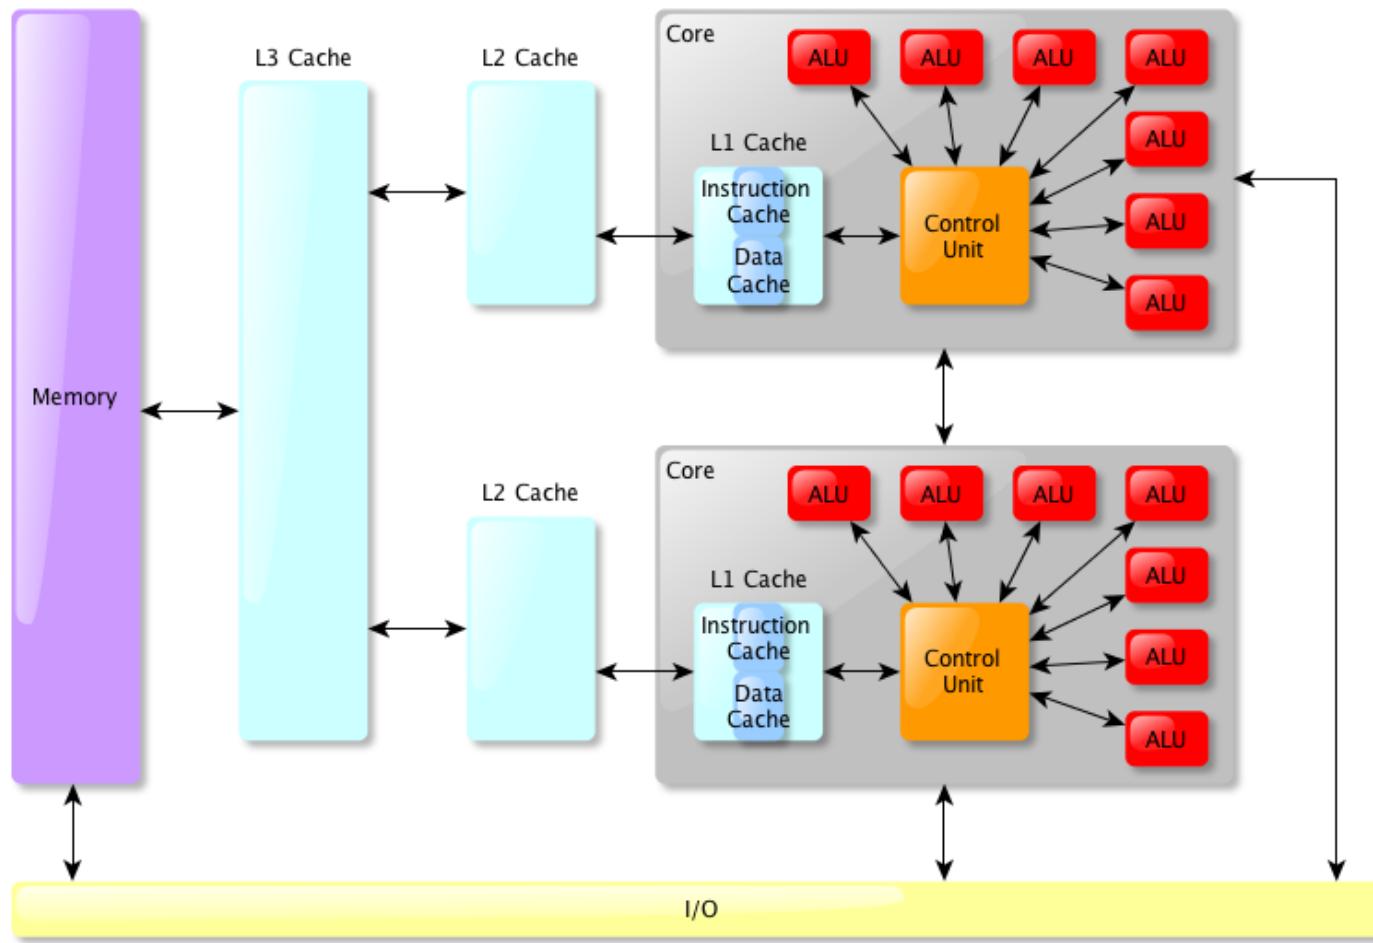


Modified Harvard Architecture

- Introduction of a caching system: smaller but much faster memory closer to the CPU (short access latency)
- Mix of both previous archs: Harvard arch to access caches; Von Neumann to access memory

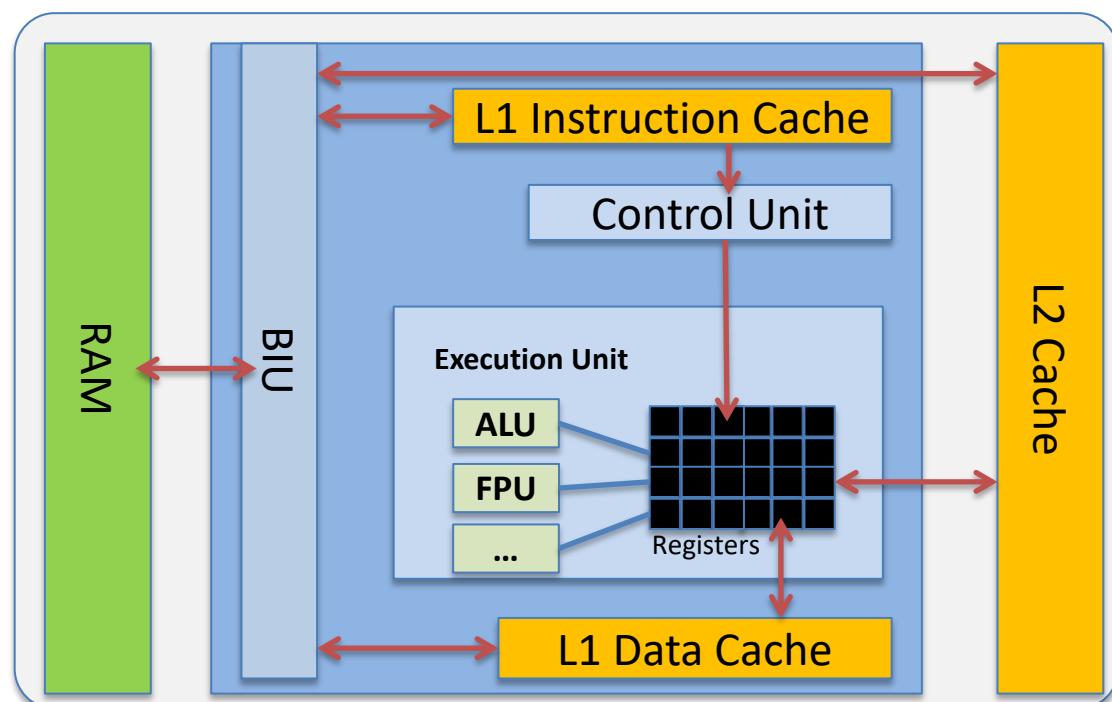


Recent architecture



Basic Concepts

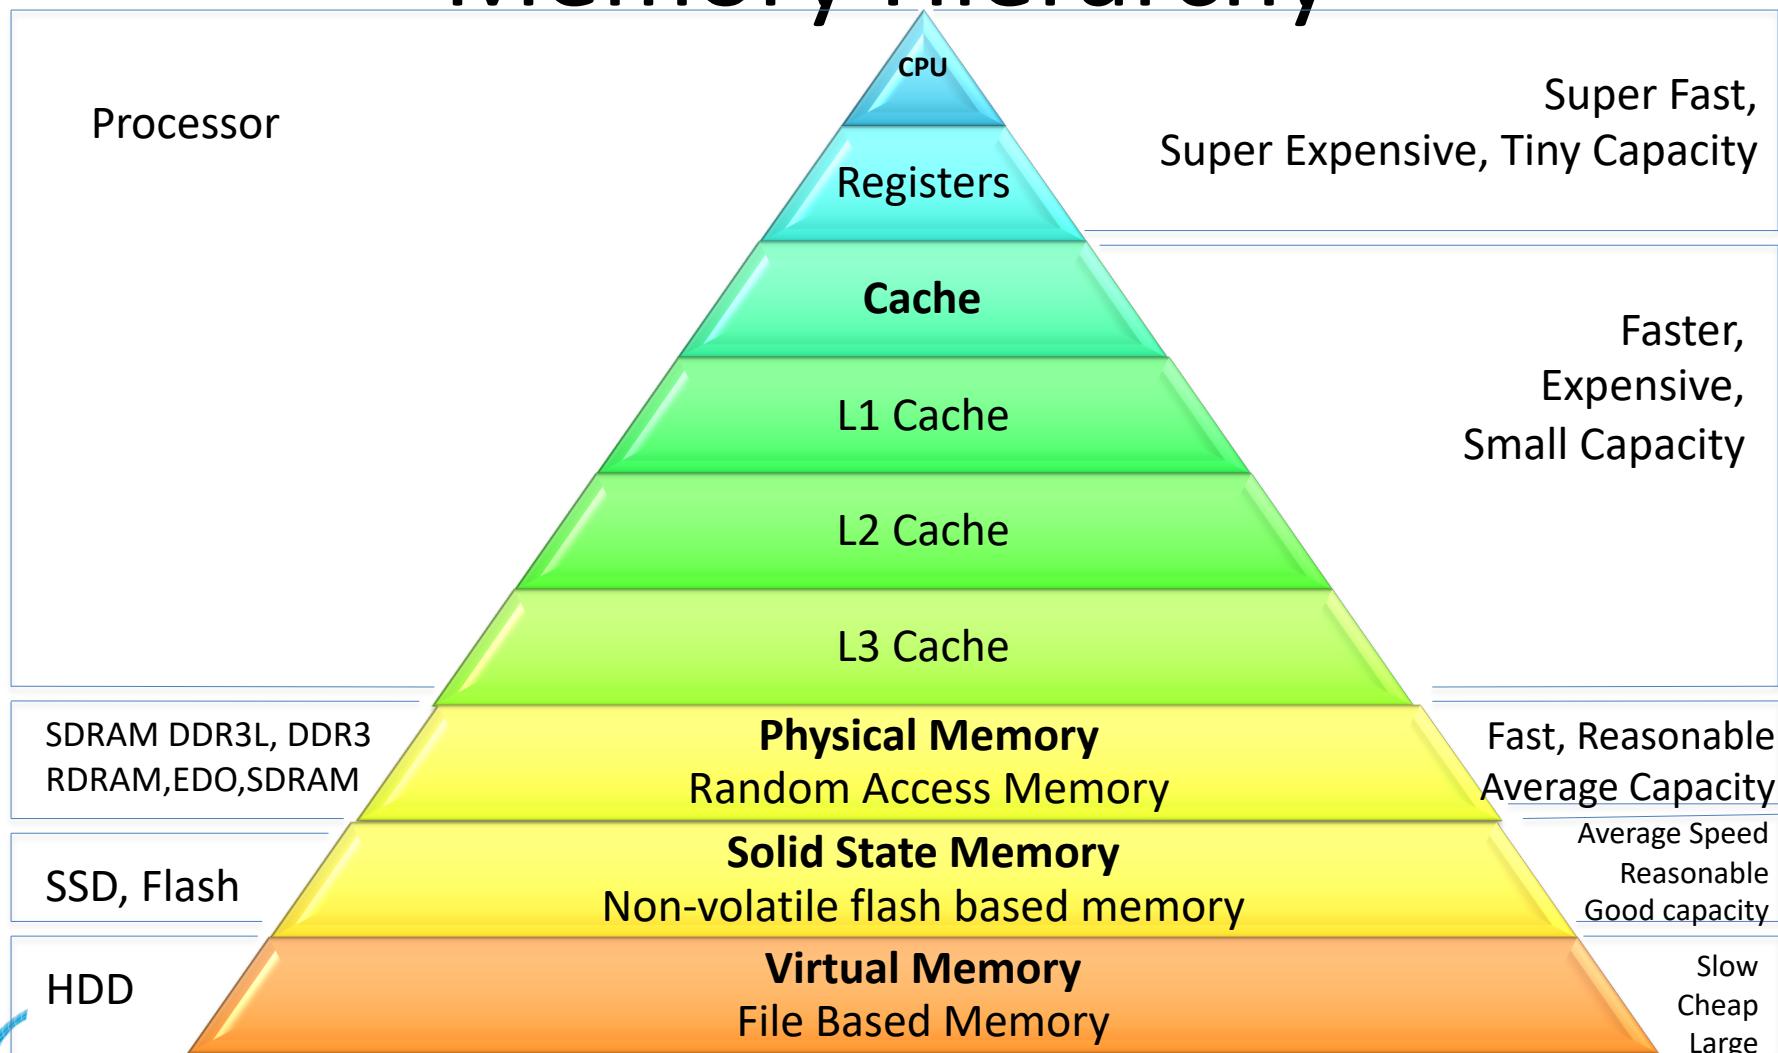
- **PROCESS**: is a task being run by a computer, often simultaneously with many other tasks each taking turns on a single central processing unit;
- **CPU** (Central Processing Unit): single- or multi-core chips that can handle multiple processes;
 - **CU**: Schedules the order of instruction execution
 - **ALU**: Performs arithmetic and logical calculations
 - **FPU**: Performs operations on floating point numbers; available in many modern processors



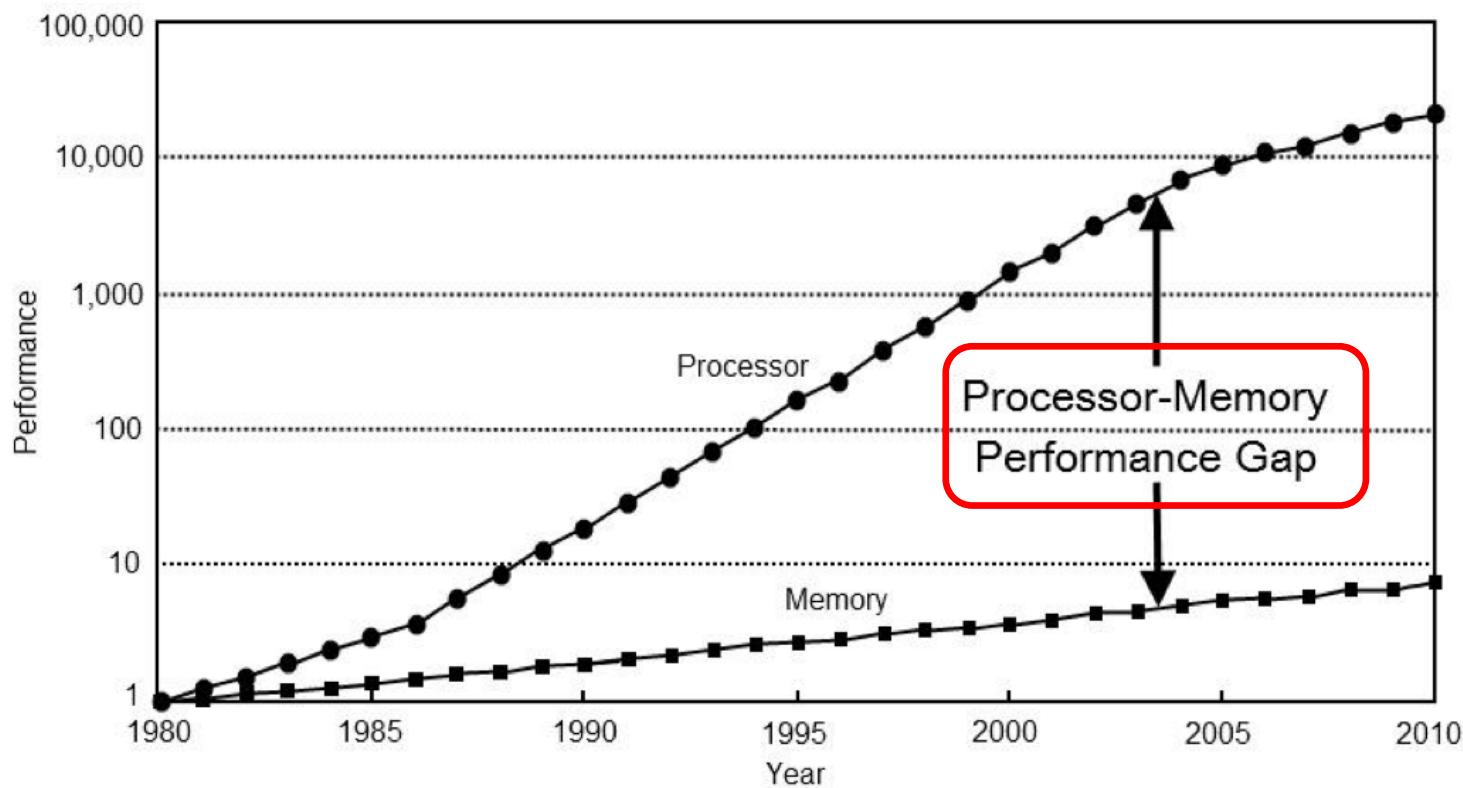
Basic Concepts

- **REGISTER:** small amount of storage where data resides that is being used right now;
- **CACHE:** small area of fast memory where data resides when it is about to be used and/or has been used recently;
- **MAIN MEMORY** (also called RAM “Random Access Memory”): where data resides when it is being used by a program that is currently running. It may be volatile or not.

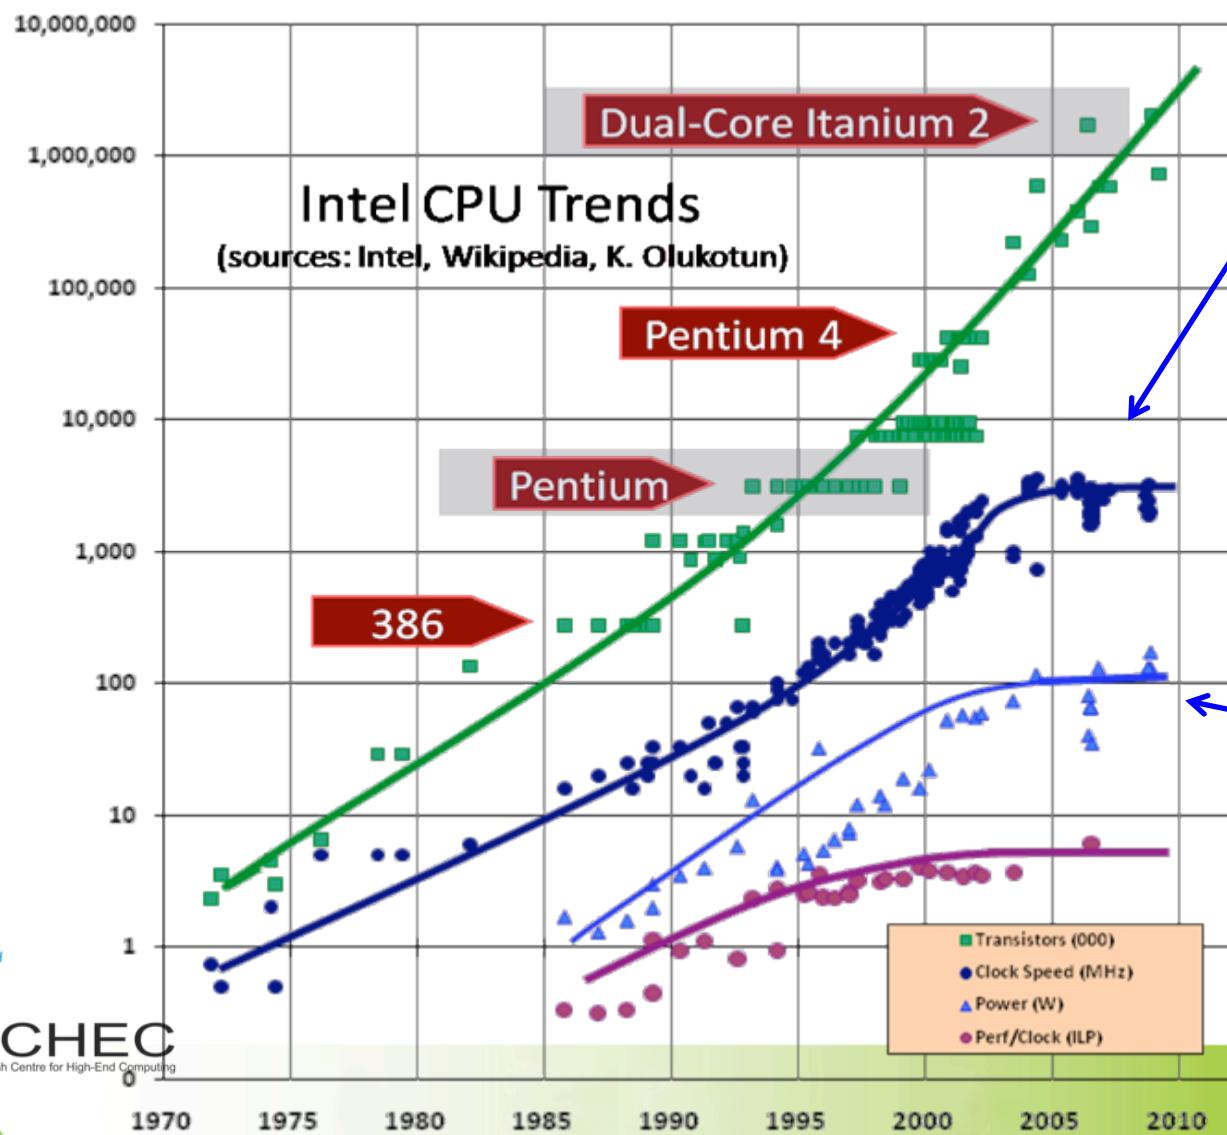
Memory Hierarchy



CPU vs Memory performance evolution



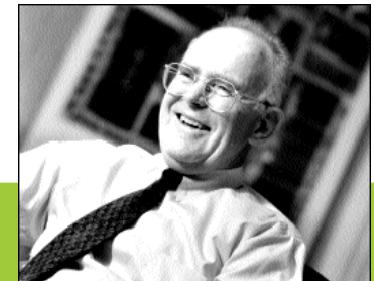
The Free Lunch is over - (Moore's Law)



“The Free Lunch
is over”:

- Processor clock-speeds are not getting any faster. Will remain ≤ 3 GHz
- Improved future performance must come from finding and exploiting *parallelism*.

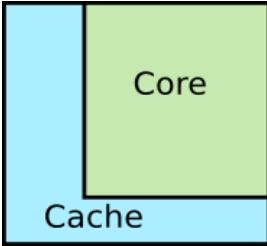
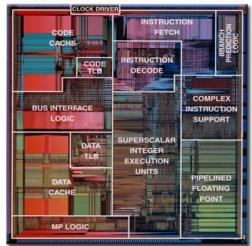
Holding down power consumption is the new imperative.



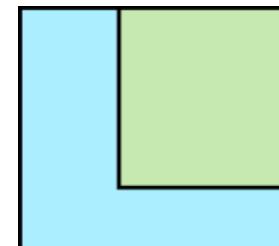
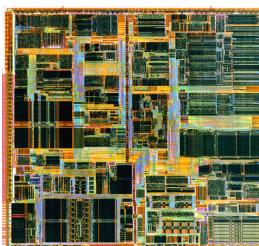
From single- to multi-core

how and why CPU evolved

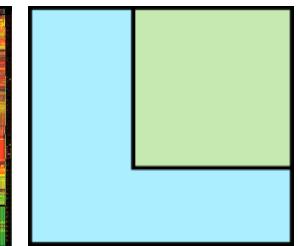
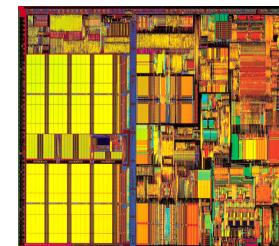
Pentium I – 1993, 60MHz



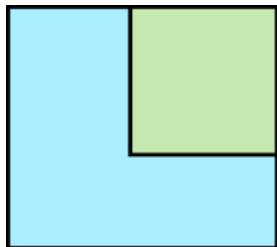
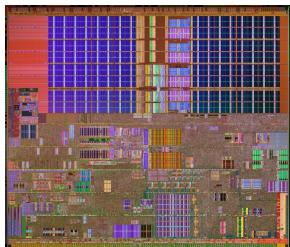
Pentium II – 1997, 233 MHz



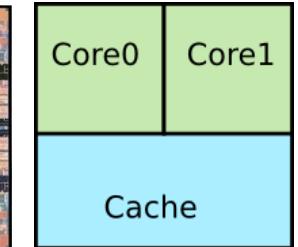
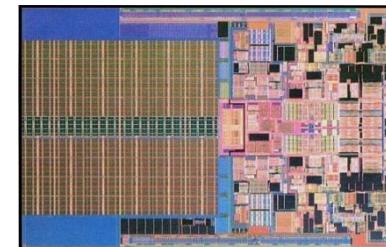
Pentium III – 1999, 500 MHz



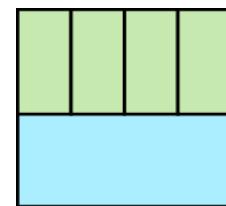
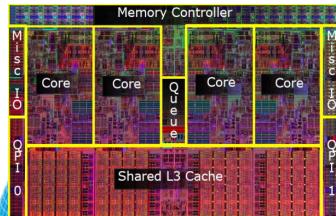
Pentium IV - 2000, 1.4 GHz



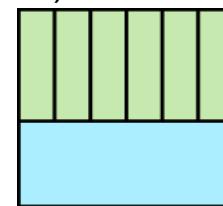
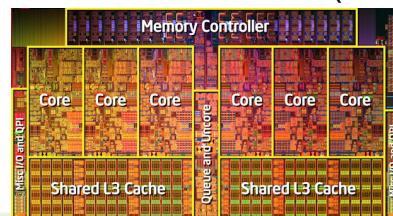
Pentium D – 2005



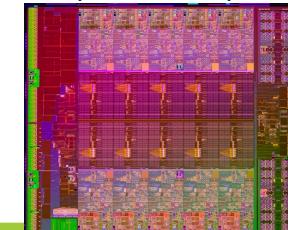
Core i7- 9xx (4 cores)



Westmere (6 cores)



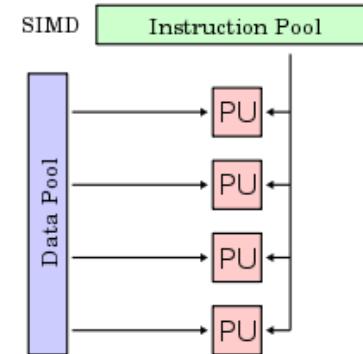
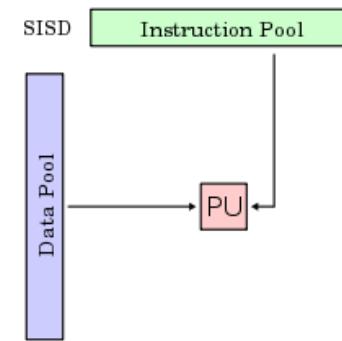
Intel Ivy Bridge-EP
(12 cores)



HPC Architectures

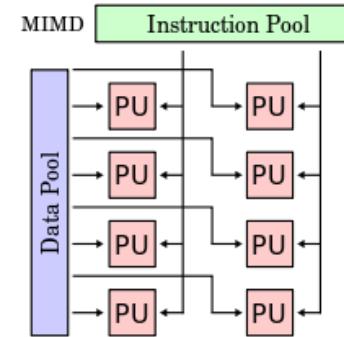
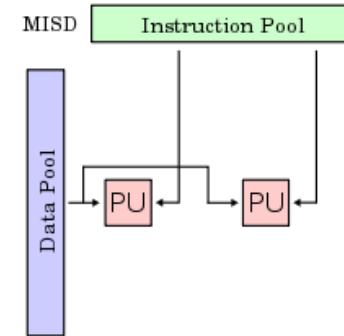
FLYNN'S TAXONOMY:

- Single Instruction Single Data (SISD): at any one time only a single instruction is executed, operating on a single data item.
- Single Instruction Multiple Data (SIMD): multiple processors, each operating on its own data item, but they are all executing the same instruction on that data item

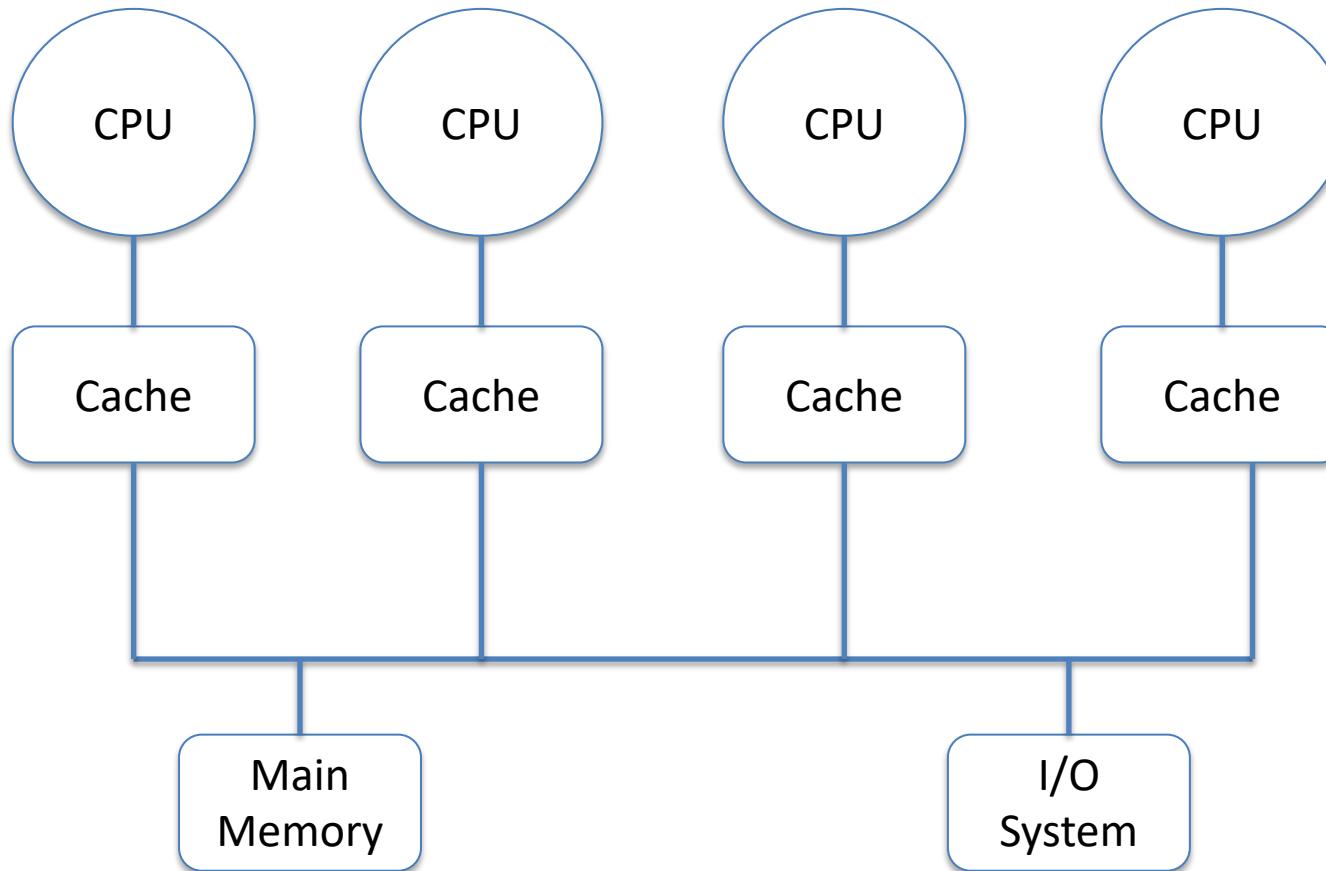


HPC Architectures

- Multiple Instructions Single Data (MISD): All processors do different operations on the same data
- Multiple Instructions Multiple Data (MIMD): multiple CPUs operate on multiple data items, each executing independent instructions
 - Shared Memory
 - Distributed Memory



Shared Memory



Shared Memory

- **UMA (Uniform Memory Access), SMP (Symmetric Multi-Processing)**
 - The latency to access any address in the logical memory space is the same for each CPU.
- **NUMA (Non-Uniform Memory Access)**
 - The latency to access any address in the logical memory space is determined by the physical distance from the CPU.
- **Cache-Coherency (cc)**
 - To ensure cache consistency (i.e. local cache has the most up-to-date copy of a shared memory resource), cache-coherency protocols are implemented on modern systems.

Shared Memory

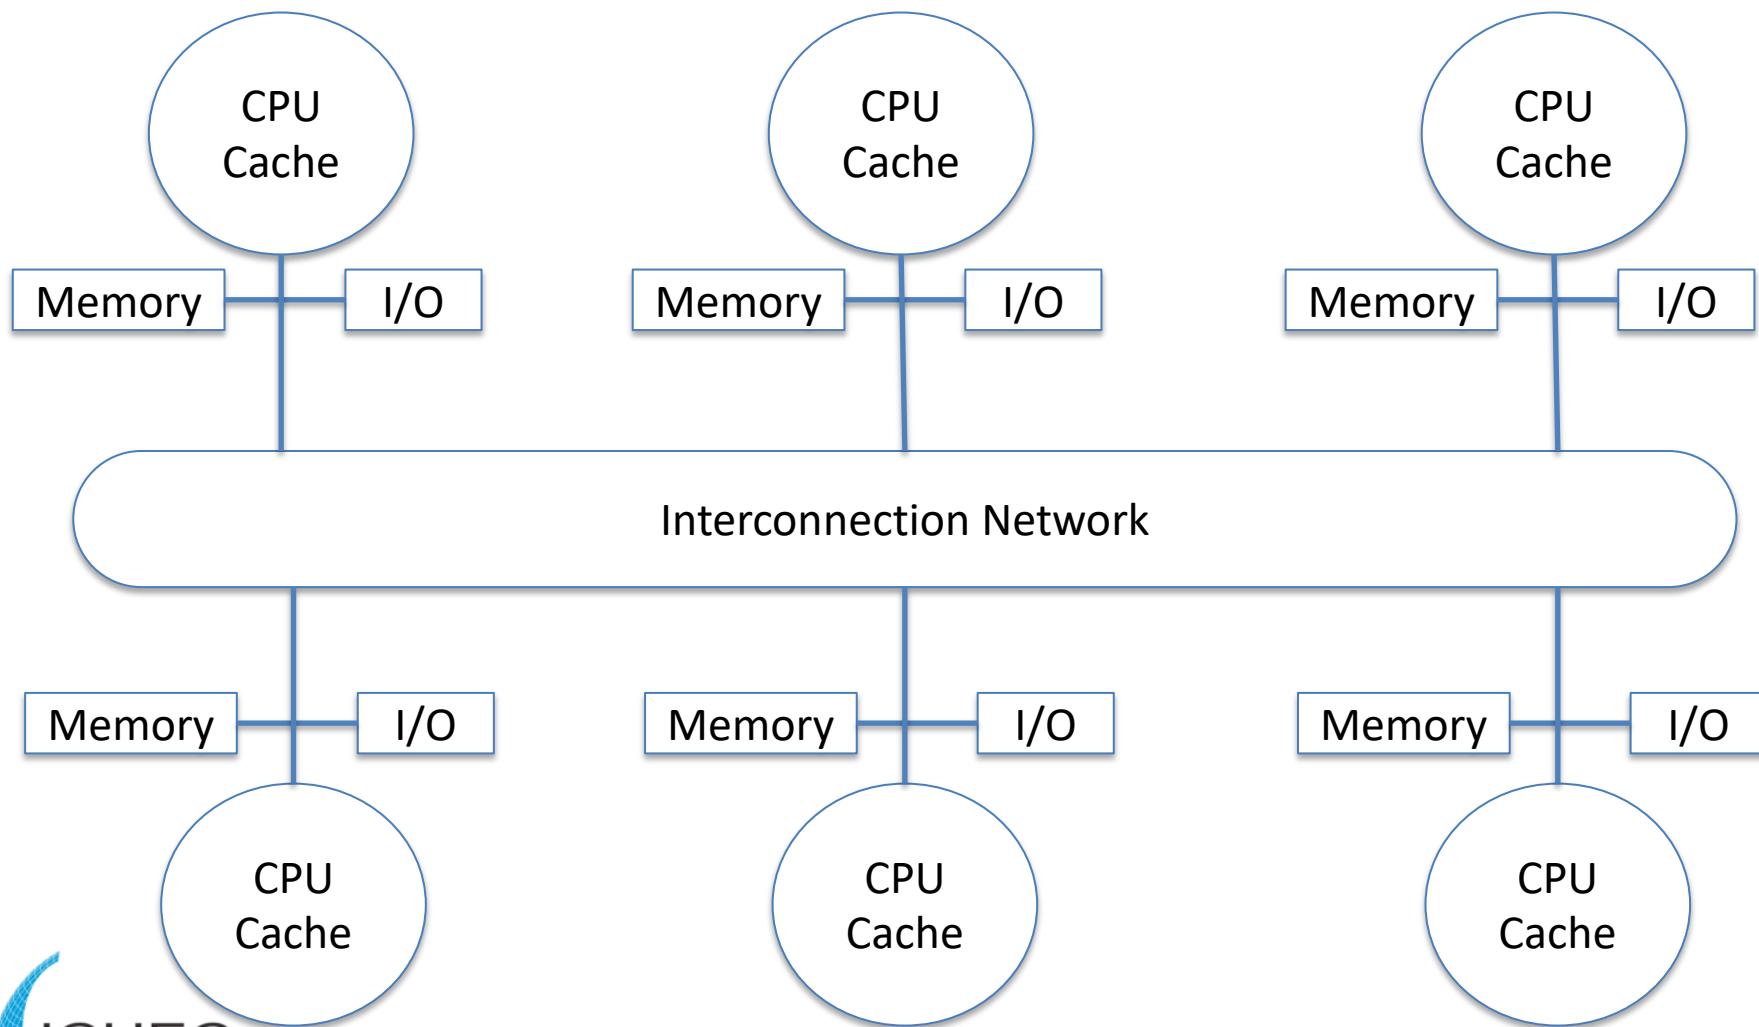
ADVANTAGES:

- Global address space provides a user-friendly programming perspective to memory (such as with the OpenMP API)
- Data sharing between tasks is both fast and uniform due to the proximity of memory to the CPUs

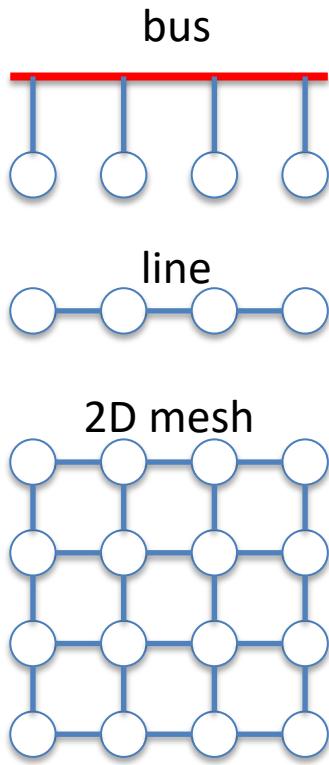
DISADVANTAGES:

- Need for cache-coherency
- Lack of scalability between memory and CPUs. Adding more CPUs can geometrically increase traffic on the shared memory-CPU path, and for cache coherent systems, geometrically increase traffic associated with cache/memory management.
- Programmer responsibility for synchronization constructs that ensure "correct" access of global memory (i.e. prevent race conditions)

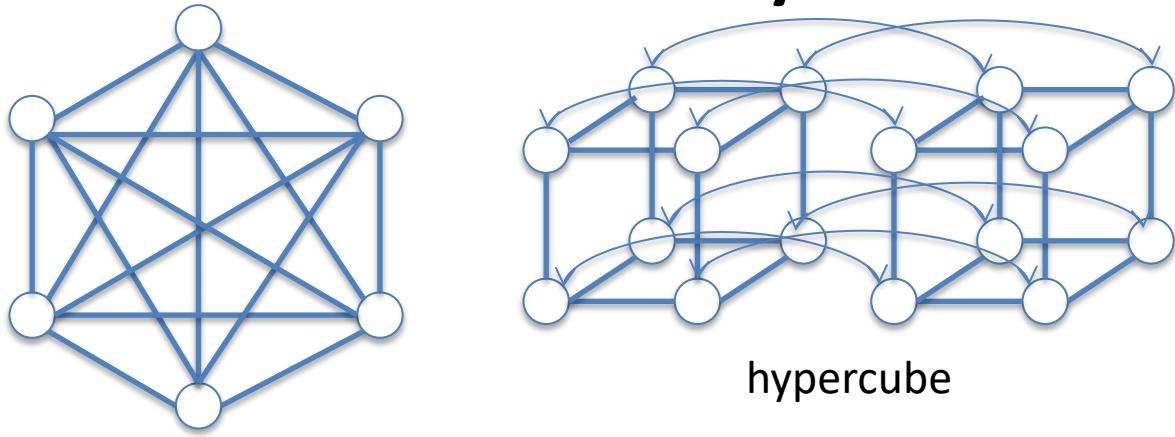
Distributed Memory



Distributed Memory



The choice of the interconnection has a **strong** impacts of the communication performance



Topology	Degree	Diameter	Avg. dist.	Bisection	Tot. BW (links)
bus	1	1	1	1	1
line	2	N-1	N/2	1	N-1
2D mesh	4	$2*(N^{1/2}-1)$	$2N^{1/2}/3$	N	$2N-2$
hypercube 2-cube	$\log_2 N$	$\log_2 N$	$\log_2 N/2$	$N/2$	$N/(2 \log_2 N)$
fully Connected	N-1	1	1	$N^2/4$	$N(N-1)/2$

Distributed Memory

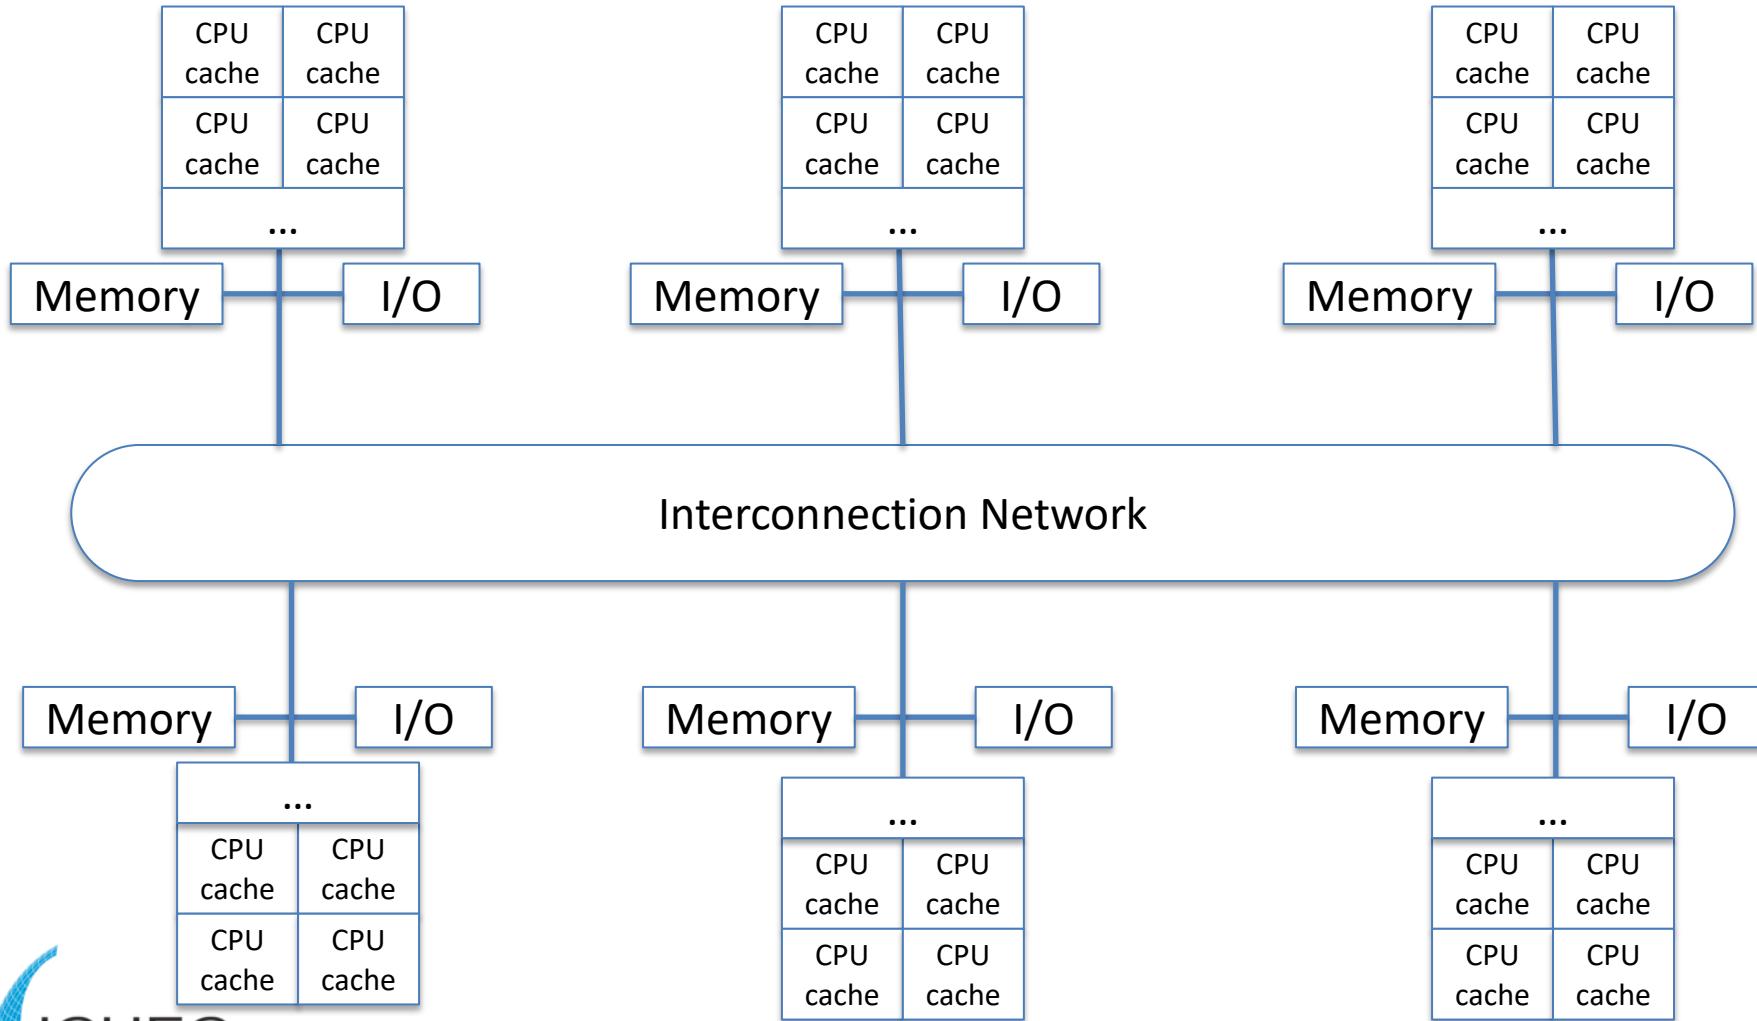
ADVANTAGES:

- Memory is scalable with number of processors. The number of processors and the size of memory increases proportionately.
- Each processor can rapidly access its own memory without interference and without the overhead incurred with trying to maintain cache coherency.
- Cost effectiveness: can use commodity, off-the-shelf processors and networking.

DISADVANTAGES:

- Programmer is responsible for mapping data structures across processors
- Programmer is responsible for coordinating communication between nodes when remote data is required in a local computation (called message-passing)
- Load balancing problem.
- Currently, only low-level programming API's (such as MPI) are available to perform message-passing

Distributed-Shared Memory



Distributed-Shared Memory

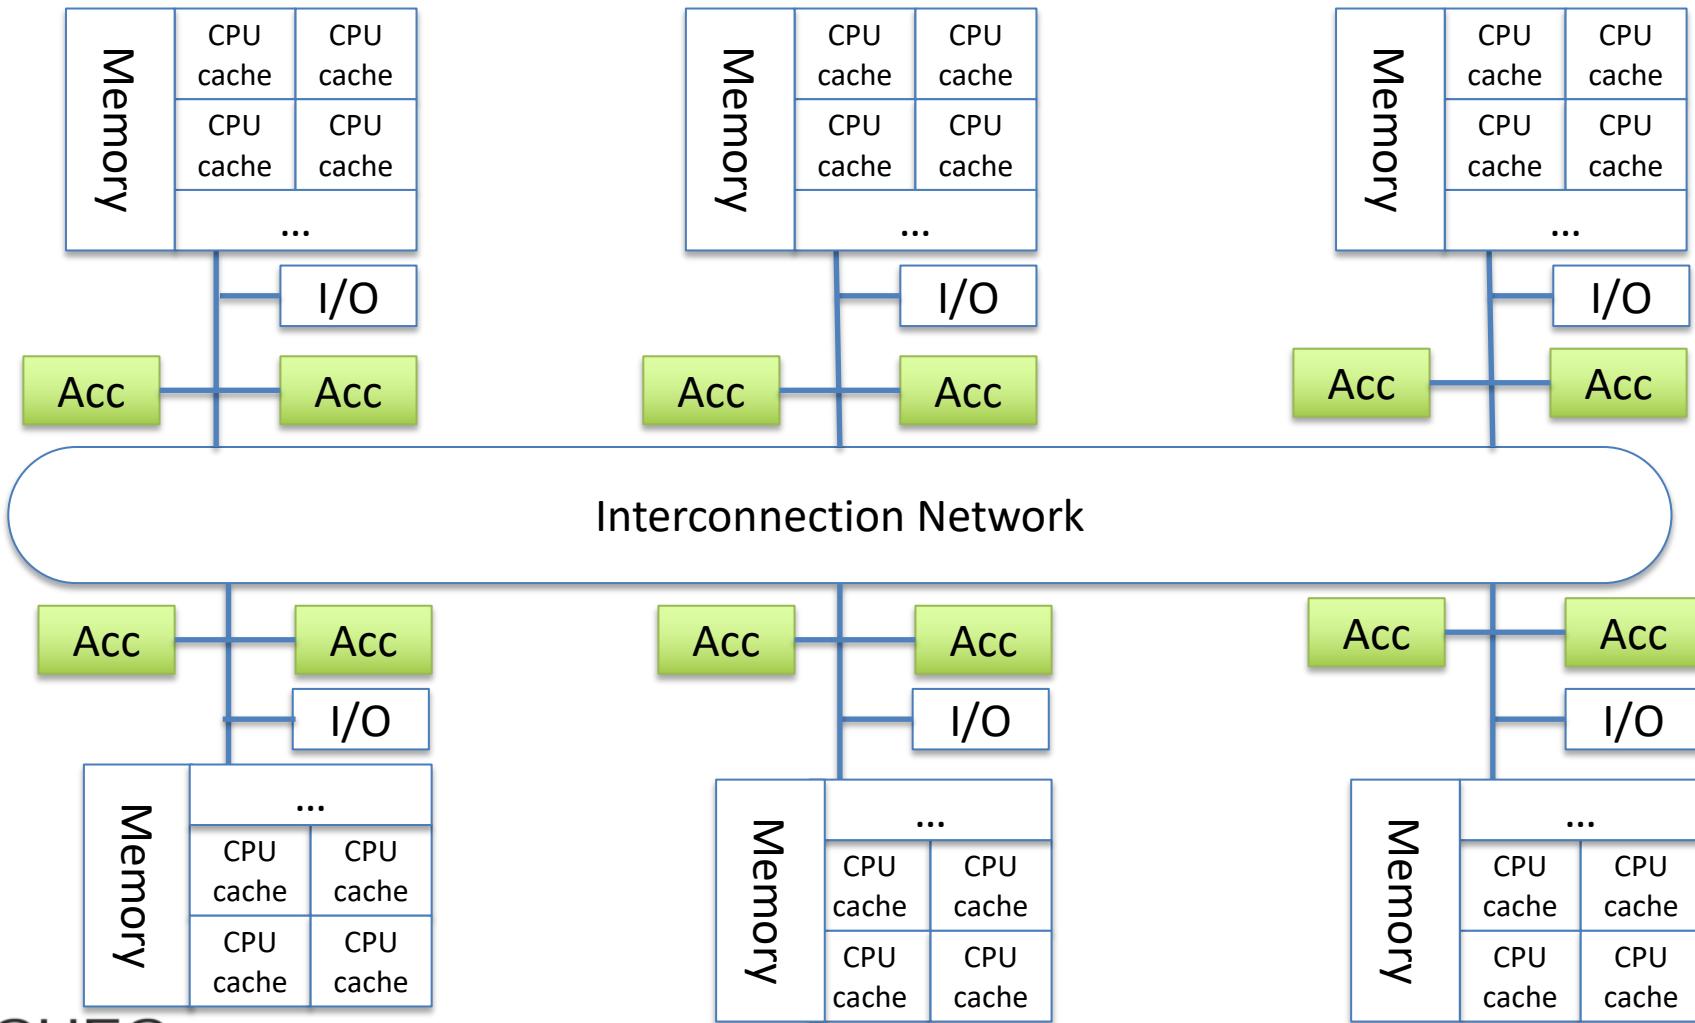
ADVANTAGES:

- Exploits the advantages of both shared and distributed-memory architectures at different levels.
- Can exploit both shared and distributed-memory programming paradigms (OpenMP and MPI) to solve difficult tasks

DISADVANTAGES:

- The same disadvantages of a pure Shared Memory system
- The same disadvantages of a simple Distributed system

Heterogeneous Systems



Heterogeneous Systems

ADVANTAGES:

- Accelerators can speed-up the calculation

DISADVANTAGES:

- Accelerators are separated from the system (like a PCI card), they do not share the memory with the host system;
- Accelerators require their own programming environment.