



# OpenMP

## Loop Level Parallelism





# Outline

- Work-sharing Constructs
  - loop
- Combined Constructs
- Reduction Clause
- Schedule Clause
- Synchronisation Constructs
  - master, critical, barrier, atomic, flush, ordered
- Monte Carlo Method
- Timing with OpenMP





# Loop Construct

- splits up loop iterations among the threads in a team.
- don't create a team of threads;
- there is an implied barrier at the end (unless `nowait` is used).

## C/C++:

```
#pragma omp parallel [clauses]  
{  
    #pragma omp for [clauses] [nowait]  
    {  
        ...  
    }  
}
```

## Fortran:

```
!$omp parallel [clauses]  
    !$omp do [clauses]  
    ...  
    !$omp end do [nowait]  
!$omp end  
parallel
```



## How is OpenMP Typically Used?

- Find your most time consuming loops.
- Split them up between threads.

### Split-up this loop between multiple threads


```
void main()
{
    double Res[1000];

    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

Sequential program

```
void main()
{
    double Res[1000];
    #pragma omp parallel
    #pragma omp for
    for(int i=0;i<1000;i++) {
        do_huge_comp(Res[i]);
    }
}
```

Parallel program





## Performance Considerations:

- Loop dependency
- Cache miss

```
for(i=1; i<n; i++)  
  for(j=1; j<n; j++)  
    a[i][j]=2*a[i-1][j];
```



```
#pragma omp parallel  
#pragma omp for private(i)  
for(int j=0; j<n; j++)  
  for(i=1; i<n; i++)  
    a[i][j]=2*a[i-1][j];
```





## Example: (Parallel matrix vector product $y=Ax$ )

**C:**

```
#pragma omp parallel
default(none), shared(n, A, x, y),
private(i, j, sum, tid)
{
tid=omp_get_thread_num();
#pragma omp for
for (i = 0; i < n; i++) {
    sum=0;
    for (j=0; j<n; j++)
        sum+=A[i*n+j]*x[j];
    y[i]=sum;
    printf("%d: y[%d]=%f\n", tid,
i, y[i]);
}
}
```

**Fortran:**

```
!$omp parallel default(none),
shared(A, x, y, n), private(i, j,
sum, tid)
tid=omp_get_thread_num()
!$omp do
do i=1,n
    sum=0
    do j=1,n
        sum=sum+A((i-1)*n+j)*x(j)
    enddo
    y(i)=sum
    print *, tid, y(i)
enddo
!$omp end do
!$omp end parallel
```





```
for(i=0; i<n*n; i++){  
    A[i]=i;  
}  
for(i=0; i<n; i++){  
    x[i]=1;  
}
```

## **Compile:** (Intel)

```
>icc -openmp hello.c -o a.out  
>ifort -openmp hello.f90 -o a.out
```

## **Execute:**

```
>export OMP_NUM_THREADS=4  
>./a.out
```

```
Thread 0: y[0]=120.000000  
Thread 0: y[1]=376.000000  
Thread 0: y[2]=632.000000  
Thread 0: y[3]=888.000000  
Thread 1: y[4]=1144.000000  
Thread 1: y[5]=1400.000000  
Thread 1: y[6]=1656.000000  
Thread 1: y[7]=1912.000000  
Thread 2: y[8]=2168.000000  
Thread 2: y[9]=2424.000000  
Thread 2: y[10]=2680.000000  
Thread 2: y[11]=2936.000000  
Thread 3: y[12]=3192.000000  
Thread 3: y[13]=3448.000000  
Thread 3: y[14]=3704.000000  
Thread 3: y[15]=3960.000000
```





# Combined Constructs

- The following shortcuts are supported:
  - parallel for / parallel do
  - parallel sections
  - parallel workshare
- Equivalent to a parallel construct followed by a work-sharing construct.

**#pragma omp parallel for**


**Same as**

**#pragma omp parallel**  
**#pragma omp for**

**!\$omp parallel do**

**Same as**

**!\$omp parallel**  
**!\$omp do**







## Parallel Loop Directives:

### C/C++:

```
#pragma omp parallel for [clause[clause]...]
```

```
for ( index = first; index <= last ; index++ ){  
    body of the loop  
}
```

### Fortran:

```
!$omp parallel do [clause[clause]...]
```

```
do index = first, last [, stride]  
    body of the loop  
enddo
```

```
!$omp end parallel do
```





# Reduction Clause

- Performs a reduction operation on the variables in the list  
C: `reduction(operator: list)`  
Fortran: `reduction(operator/intrinsic: list)`
- A local copy of each list variable is created
- Compiler finds the reduction operation and updates local copy
- Combine local copies into a single variable
- for, parallel, sections
- Reduction operations: `+, -, *, ...` (C/C++/Fortran)  
`&&, ||, ...` (C/C++) `.AND., .OR., ...` (Fortran)  
`max, min, ...` (Fortran)



## Example: (Parallel matrix vector product $y=Ax$ )

```
C:
for (i = 0; i < n; i++) {
    sum=0;
    #pragma omp parallel for
    default(shared) private(j, tid)
    reduction(+:sum)
    for (j=0; j<n; j++){
        sum=sum+(A[i*n+j]*x[j]);
        if(i==0){
            tid=omp_get_thread_num();
            printf("Tid=%d, A[%d]=%f,
sum=%f\n", tid, i*n+j, A[i*n+j], sum);
        }
    }
    y[i]=sum;
    printf("y[%d]=%f\n", i, y[i]);
}
```

>./a.out

Tid=0, A[0]=0.000000, sum=0.000000  
Tid=0, A[1]=1.000000, sum=1.000000  
Tid=0, A[2]=2.000000, sum=3.000000  
Tid=0, A[3]=3.000000, sum=6.000000  
Tid=1, A[4]=4.000000, sum=4.000000  
Tid=1, A[5]=5.000000, sum=9.000000  
Tid=1, A[6]=6.000000, sum=15.000000  
Tid=1, A[7]=7.000000, sum=22.000000  
Tid=2, A[8]=8.000000, sum=8.000000  
Tid=2, A[9]=9.000000, sum=17.000000  
Tid=2, A[10]=10.000000, sum=27.000000  
Tid=2, A[11]=11.000000, sum=38.000000  
Tid=3, A[12]=12.000000, sum=12.000000  
Tid=3, A[13]=13.000000, sum=25.000000  
Tid=3, A[14]=14.000000, sum=39.000000  
Tid=3, A[15]=15.000000, sum=54.000000





# Schedule Clause

- Describes how iterations of the loop are divided among the team threads

C: `schedule (type [,chunk])`

Fortran: `schedule (type [,chunk])`


- Types:
  - `schedule(static [,chunk])`
  - `schedule(dynamic [,chunk])`
  - `schedule(guided [,chunk])`
  - `schedule(runtime)`

- More: **ordered** clause





## SCHEDULE(STATIC):



- Iterations are divided evenly among threads
- Divides the work into chunk sized parcels
- If there are N threads, each does every N<sup>th</sup> chunk of work

### Fortran:

```
!$omp parallel do schedule(static,3)
```

```
do i = 1, 36
  Work (i)
end do
```


```
!$omp end do
```

OpenMP Loop Level Parallelism





## SCHEDULE(DYNAMIC):



- Divides the workload into chunk sized parcels
- As a thread finishes one chunk, it grabs the next available chunk
- More overhead, but better load balancing

### Fortran:

```
!$omp parallel do schedule(dynamic,1)
```

```
do i = 1, 36
```

```
  Work (i)
```

```
end do
```


```
!$omp end do
```

OpenMP Loop Level Parallelism





## SCHEDULE(GUIDED):



- Iterations are divided into chunks such that the size of each successive chunk decreases.
- chunk: the size of the smallest chunk size

### Fortran:

```
!$omp parallel do schedule(guided,1)
```

```
do i = 1, 36
  Work (i)
end do
```

```
!$omp end do
```

OpenMP Loop Level Parallelism







# Synchronisation Constructs

- To control how the execution of each thread proceeds relative to other team threads
- To protect against data races
- Synchronisation Constructs:
  - master
  - critical
  - barrier
  - atomic
  - flush
  - ordered







## Master directive:

- Only **master** thread can enter the structured block

**C/C++:**

```
#pragma omp master  
...
```

**Fortran:**

```
!$omp master  
...  
!$omp end  
  
master
```

## Critical directive:

- Only one thread at a time can enter a **critical** section

**C/C++:**

```
#pragma omp critical [name]  
...
```

**Fortran:**

```
!$omp critical [name]  
...  
!$omp end  
  
critical
```





## Barrier directive:

- Threads wait each other until all have reached that **barrier**

**C/C++:**

```
#pragma omp barrier
```

**Fortran:**

```
!$omp barrier
```

## Atomic directive:

- Access to a specific memory location **atomically**
- $x \text{ binop} = \text{expr}$ ,  $x++$ ,  $--x$ , ... (C/C++)
- $x = x \text{ op expr}$ ,  $x = \text{expr op } x$ , ... (Fortran)

**C/C++:**

```
#pragma omp atomic  
statement
```

**Fortran:**

```
!$omp atomic  
statement
```





## Flush directive:

- A synchronization point at which thread visible variables are **flushed**

**C/C++:**

```
#pragma omp flush (list)
```

**Fortran:**

```
!$omp flush  
(list)
```

## Ordered directive:

- Iterations of the enclosed loop will be executed in the same **order** as if they were executed sequentially

**C/C++:**

```
#pragma omp ordered  
...
```

**Fortran:**

```
!$omp ordered  
...  
!$omp end ordered
```






# Monte Carlo Method

- Based on the use of random numbers and probability statistics
- Applied in many research areas
  - Financial Analysis
  - Nuclear Physics
  - Population Dynamics
  - Weather Prediction
- Random Number Generation
- Very easy to parallelize





## Calculation of pi:



- Area of the shaded circle:  $\pi r^2/4$
- Area of the square:  $r^2$
- Probability  $= \pi/4$
  
- $r=1$  at  $(0,0)$
- distance from the origin:  $\sqrt{x^2+y^2}$
- If the distance  $\leq 1$ , then it is a hit.

- $\Pi = 4 \times \text{total number of hits} / \text{total number of random numbers}$



C:

```
#pragma omp parallel for default(shared), private(i),  
reduction(+:hit), schedule(dynamic, 1)  
for (i = 0; i < rnums; i++) {  
    double x = (double) rand() / RAND_MAX;  
    double y = (double) rand() / RAND_MAX;  
    if (x * x + y * y <= 1)  
        hit++;  
}  
printf("The number of hits=%d, Estimate of pi = %f\n", hit, 4.0  
* hit / rnums);
```

True value of pi = 3.141593

rnums=10: The number of hits=8, Estimate of pi = 3.200000

rnums=100: The number of hits=78, Estimate of pi = 3.120000

rnums=1000: The number of hits=783, Estimate of pi = 3.132000

rnums=10000: The number of hits=7909, Estimate of pi = 3.163600

rnums=100000: The number of hits=78524, Estimate of pi = 3.140960



C:

```
#pragma omp parallel default(shared), private(thit, i, tid)
{
thit=0;
#pragma omp for schedule(dynamic, 1)
for (i = 0; i < rnums; i++) {
    double x = (double) rand() / RAND_MAX;
    double y = (double) rand() / RAND_MAX;
    if (x * x + y * y <= 1)
        thit++;
}
tid=omp_get_thread_num();
printf("Tid=%d thit=%d\n", tid, thit);
#pragma omp critical
    hit+=thit;
}
```

Tid=0 thit=22484, Tid=3 thit=20419, Tid=2 thit=18034, Tid=1  
thit=17728. The number of hits=78665, Estimate of pi = 3.146600





# Timing with OpenMP

- `omp_get_wtime`  
C: `double omp_get_wtime(void)`  
Fortran: `double precision function omp_get_wtime()`
- Returns a double-precision floating point value equal to the number of elapsed seconds since some point in the past.

```
double start, end, work;  
start = omp_get_wtime();  
... work to be timed ...  
end = omp_get_wtime();  
work=end-start;
```

- `omp_get_wtick`

