



Practical3

ICHEC

April 26, 2020

1 Hello world!

Build the MPI hello world example from Fig. 1.

1. Run it using one process and using 12 and 24 processes. Can you rely on the output order? Use `(mpirun -n # ./prog)` to run a MPI enabled program. The `(-n)` option takes a number, which is the number of MPI processes.
2. Change the code so it prints a message before `MPI_Init` call. When you run it using 12 processes, how many time is this message printed.
3. Change the code so that only one process handles the output. Choose yourself the right process for that.
4. Can you make the original code to print the messages in a deterministic order (ordered by rank)?

```
#include <stdio.h>
#include <mpi.h>

int main(int argc, char **argv){

    int ierror;
    int myRank, uniSize;
    int version, subversion;
    int iMyName;
    char myName[MPI_MAX_PROCESSOR_NAME];

    ierror=MPI_Init(&argc,&argv);
    ierror=MPI_Comm_size(MPI_COMM_WORLD,&uniSize);
    ierror=MPI_Comm_rank(MPI_COMM_WORLD,&myRank);
    ierror=MPI_Get_processor_name(myName,&iMyName);
    ierror=MPI_Get_version(&version,&subversion);
    printf("I am process %d out of %d running on %s with MPI version %d.%d\n",
           myRank, uniSize, myName, version, subversion);
    ierror=MPI_Finalize();
    return ierror;
}

program helloMPI
    use mpi
    implicit none

    integer :: ierror
    integer :: myRank, uniSize, version, subversion
    character(len=MPI_MAX_PROCESSOR_NAME) :: myName
    integer :: status(MPI_STATUS_SIZE), iMyName

    call MPI_Init(ierror)
    call MPI_Comm_size(MPI_COMM_WORLD, uniSize, ierror)
    call MPI_Comm_rank(MPI_COMM_WORLD, myRank, ierror)
    call MPI_Get_processor_name(myName, iMyName, ierror)
    call MPI_Get_version(version, subversion, ierror)
    write(*, '(a,i0,a,i0,a,a,a,i0,a,i0)') "I am process ", myRank, &
        " out of ", uniSize, " running on ", trim(myName), &
        " with MPI version ", version, ".", subversion
    call MPI_Finalize(ierror)
end program helloMPI
```

Figure 1. MPI hello world, top C and bottom Fortran

2 Ping Pong

Write a simple ping pong program. This involves passing of a message between two processes P_0 and P_1 . The algorithm is as follows

1. let us assume we have an initial message which contains the integer value 10.
2. P_0 increments this message by one and passes it to P_1 (ping)

3. P_1 receives the message it increments it by one and passes it back to P_0
4. the last two steps are repeated n times. (see fig. 2)

The code shall do

1. rank 0 prints the value of the message it has after n exchanges.
2. rank 0 prints the average time per exchange. **Hint:** use `MPI_Wtime()` function to get the time.
3. determine the value n for which the measured time is meaningful. **Hint:** check the resolution of the timer with `MPI_Wtick()`
4. use `MPI_Send` and `MPI_Recv` to pass the messages around.




Figure 2. Ping Pong algorithm

3 Latency and Bandwidth

Latency is defined as the time to transfer a zero length message. Bandwidth is defined as the size of the message in bytes/ transfer time.

1. Modify the ping pong code to measure the latency (use `MPI_BYTE` as transfer type).
2. Modify the ping pong code to measure the Bandwidth (use `double` or `real(kind=8)`). Measure the bandwidth for the following sizes 8 B, 512 B, 32 KiB, 2MiB these correspond to arrays of length 1 , 2^6 , 2^{12} and 2^{18} , respectively.
3. do the measurements for the following two cases: i) inside one node (`mpirun -ppn 2`), ii) in between two nodes (`mpirun -ppn 1`)